

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

In [2]:

```
data = pd.read_csv("room_occupancy.csv")
```

In [3]:

```
data
```

Out[3]:

	date	Temperature	Humidity	Light	CO2	HumidityRatio	Occupancy
0	02-02-2015 14:19	23.7000	26.2720	585.200000	749.200000	0.004764	1
1	02-02-2015 14:19	23.7180	26.2900	578.400000	760.400000	0.004773	1
2	02-02-2015 14:21	23.7300	26.2300	572.666667	769.666667	0.004765	1
3	02-02-2015 14:22	23.7225	26.1250	493.750000	774.750000	0.004744	1
4	02-02-2015 14:23	23.7540	26.2000	488.600000	779.000000	0.004767	1
...
10803	10-02-2015 09:29	21.0500	36.0975	433.000000	787.250000	0.005579	1
10804	10-02-2015 09:29	21.0500	35.9950	433.000000	789.500000	0.005563	1
10805	10-02-2015 09:30	21.1000	36.0950	433.000000	798.500000	0.005596	1
10806	10-02-2015 09:32	21.1000	36.2600	433.000000	820.333333	0.005621	1
10807	10-02-2015 09:33	21.1000	36.2000	447.000000	821.000000	0.005612	1

10808 rows × 7 columns

In [4]:



```
data.head()
```

Out[4]:

	date	Temperature	Humidity	Light	CO2	HumidityRatio	Occupancy
0	02-02-2015 14:19	23.7000	26.272	585.200000	749.200000	0.004764	1
1	02-02-2015 14:19	23.7180	26.290	578.400000	760.400000	0.004773	1
2	02-02-2015 14:21	23.7300	26.230	572.666667	769.666667	0.004765	1
3	02-02-2015 14:22	23.7225	26.125	493.750000	774.750000	0.004744	1
4	02-02-2015 14:23	23.7540	26.200	488.600000	779.000000	0.004767	1

In [5]:



```
data.tail()
```

Out[5]:

	date	Temperature	Humidity	Light	CO2	HumidityRatio	Occupancy
10803	10-02-2015 09:29	21.05	36.0975	433.0	787.250000	0.005579	1
10804	10-02-2015 09:29	21.05	35.9950	433.0	789.500000	0.005563	1
10805	10-02-2015 09:30	21.10	36.0950	433.0	798.500000	0.005596	1
10806	10-02-2015 09:32	21.10	36.2600	433.0	820.333333	0.005621	1
10807	10-02-2015 09:33	21.10	36.2000	447.0	821.000000	0.005612	1

In [6]:



```
data.shape
```

Out[6]:

```
(10808, 7)
```

In [7]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10808 entries, 0 to 10807
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                   10808 non-null  object
1   Temperature            10808 non-null  float64
2   Humidity               10808 non-null  float64
3   Light                  10808 non-null  float64
4   CO2                    10808 non-null  float64
5   HumidityRatio          10808 non-null  float64
6   Occupancy              10808 non-null  int64
dtypes: float64(5), int64(1), object(1)
memory usage: 591.2+ KB
```

In [8]:

```
data.describe()
```

Out[8]:

	Temperature	Humidity	Light	CO2	HumidityRatio	Occupancy
count	10808.000000	10808.000000	10808.000000	10808.000000	10808.000000	10808.000000
mean	20.819992	25.638407	137.694088	634.00507	0.003903	0.249907
std	1.078410	4.953792	212.175483	312.81727	0.000803	0.432979
min	19.000000	16.745000	0.000000	412.75000	0.002674	0.000000
25%	20.000000	21.390000	0.000000	441.00000	0.003323	0.000000
50%	20.700000	25.680000	0.000000	464.00000	0.003805	0.000000
75%	21.500000	28.324167	413.541667	761.00000	0.004372	0.000000
max	24.408333	39.117500	1697.250000	2028.50000	0.006476	1.000000

In [9]:

```
data.isnull().sum()
```

Out[9]:

```
date           0
Temperature    0
Humidity       0
Light          0
CO2            0
HumidityRatio  0
Occupancy      0
dtype: int64
```

In [10]:

```
data.describe().columns
```

Out[10]:

```
Index(['Temperature', 'Humidity', 'Light', 'CO2', 'HumidityRatio',  
      'Occupancy'],  
      dtype='object')
```

In [11]:

```
data['Occupied'] = data['Occupancy'].replace({0: 'No', 1: 'Yes'})
```

In [12]:

```
data.head()
```

Out[12]:

	date	Temperature	Humidity	Light	CO2	HumidityRatio	Occupancy	Occupied
0	02-02-2015 14:19	23.7000	26.272	585.200000	749.200000	0.004764	1	Yes
1	02-02-2015 14:19	23.7180	26.290	578.400000	760.400000	0.004773	1	Yes
2	02-02-2015 14:21	23.7300	26.230	572.666667	769.666667	0.004765	1	Yes
3	02-02-2015 14:22	23.7225	26.125	493.750000	774.750000	0.004744	1	Yes
4	02-02-2015 14:23	23.7540	26.200	488.600000	779.000000	0.004767	1	Yes

In [13]:

```
data['Occupied'].value_counts()
```

Out[13]:


```
No      8107  
Yes     2701  
Name: Occupied, dtype: int64
```

In [23]:

```
data.Occupied.value_counts().sort_index(ascending=False).plot(kind='bar')
```

Out[23]:

<matplotlib.axes._subplots.AxesSubplot at 0x5bc14ea340>



In [25]:

```
features=['Temperature', 'Humidity', 'Light', 'CO2', 'HumidityRatio']  
data.groupby('Occupied')[features].mean()
```

Out[25]:

	Temperature	Humidity	Light	CO2	HumidityRatio
Occupied					
No	20.449638	25.115111	25.595423	502.284955	0.003724
Yes	21.931603	27.209071	474.156093	1029.360484	0.004440

In [26]:

```
data_features = data[['Temperature', 'Humidity', 'Light', 'CO2', 'HumidityRatio', 'Occupied']]
```

In [27]:


```
print(data_features.corr())
sns.heatmap(data_features.corr())
```

	Temperature	Humidity	Light	CO2	HumidityRatio	\
Temperature	1.000000	-0.040270	0.685543	0.640987	0.303167	
Humidity	-0.040270	1.000000	0.104641	0.469475	0.938169	
Light	0.685543	0.104641	1.000000	0.693588	0.335745	
CO2	0.640987	0.469475	0.693588	1.000000	0.686438	
HumidityRatio	0.303167	0.938169	0.335745	0.686438	1.000000	
Occupancy	0.595005	0.183020	0.915363	0.729540	0.386346	

	Occupancy
Temperature	0.595005
Humidity	0.183020
Light	0.915363
CO2	0.729540
HumidityRatio	0.386346
Occupancy	1.000000

Out[27]:


<matplotlib.axes._subplots.AxesSubplot at 0x5bc16af640>



In [28]:




```
sns.displot(kind='kde', x='Light', hue='Occupied', data=data, aspect=2);
```



In [29]:



```
sns.displot(kind='kde', x='CO2', hue='Occupied', data=data, aspect=2);
```

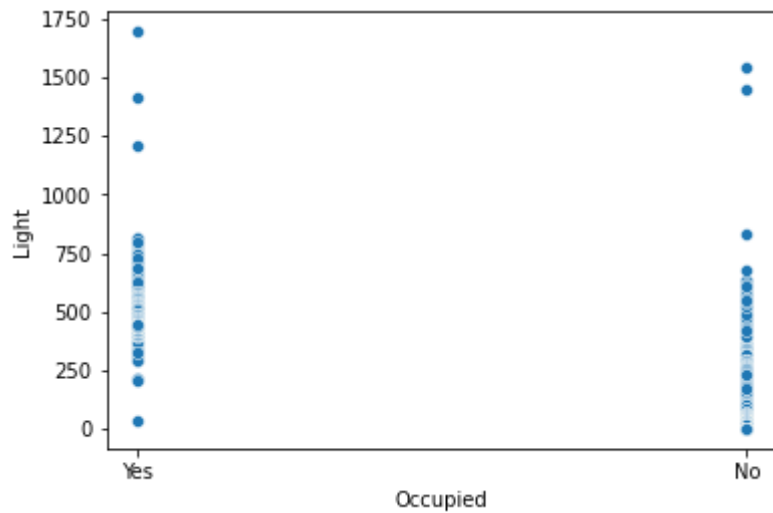


In [31]:

```
sns.scatterplot(y="Light",x="Occupied",data=data)
```

Out[31]:

<matplotlib.axes._subplots.AxesSubplot at 0x5ba6cb8250>

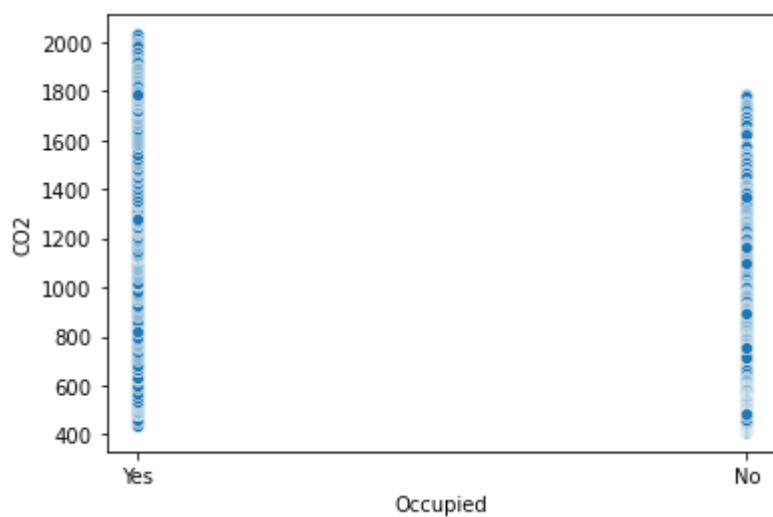


In [32]:

```
sns.scatterplot(y="CO2",x="Occupied",data=data)
```


Out[32]:

<matplotlib.axes._subplots.AxesSubplot at 0x5b926b2d30>




In [34]:

```
sns.boxplot(x='Light', y='Occupied', data=data);
```



In [35]:

```
sns.boxplot(x='CO2', y='Occupied', data=data);
```



In [37]:

```
x = data.drop(['date', 'Occupancy', 'Occupied'], axis = 1)
```

In [38]:

```
y = data.Occupancy
```

In [39]:

```
x.shape
```

Out[39]:

```
(10808, 5)
```

In [40]:

```
y.shape
```

Out[40]:

```
(10808,)
```

In [49]:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)
```

In [50]:

```
model = DecisionTreeClassifier()
model.fit(X_train, y_train)
```

Out[50]:

```
DecisionTreeClassifier()
```

In [51]:

```
y_pred = model.predict(X_test)
```

In [52]:

```
print("Training Accuracy :", model.score(X_train, y_train))
print("Testing Accuracy :", model.score(X_test, y_test))
```

```
Training Accuracy : 1.0
Testing Accuracy : 0.9861239592969473
```

In [53]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)
```

In [54]:

```
classifier= RandomForestClassifier(n_estimators= 10, criterion="entropy")
classifier.fit(X_train, y_train)
```

Out[54]:

```
RandomForestClassifier(criterion='entropy', n_estimators=10)
```

In [55]:

```
y_pred = classifier.predict(X_test)
```

In [56]:

```
print("Training Accuracy :", classifier.score(X_train, y_train))
print("Testing Accuracy :", classifier.score(X_test, y_test))
```

```
Training Accuracy : 0.9987277353689568
```

```
Testing Accuracy : 0.9939870490286772
```

In [57]:

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)
```

In [58]:

```
model1= KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2)
model1.fit(X_train, y_train)
```

Out[58]:

```
KNeighborsClassifier()
```

In [59]:

```
y_pred = model1.predict(X_test)
```

In [60]:

```
print("Training Accuracy :", model1.score(X_train, y_train))
print("Testing Accuracy :", model1.score(X_test, y_test))
```

```
Training Accuracy : 0.9899375433726578
```

```
Testing Accuracy : 0.9893617021276596
```

In [61]:

```
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)
```

In [62]:



```
model2 = GaussianNB()  
model2.fit(X_train, y_train)
```

Out[62]:

GaussianNB()

In [63]:



```
y_pred = model2.predict(X_test)
```

In [64]:



```
print("Training Accuracy :", model2.score(X_train, y_train))  
print("Testing Accuracy :", model2.score(X_test, y_test))
```

Training Accuracy : 0.9686560259079343

Testing Accuracy : 0.96577243293247