

Understanding the Gender Pay Gap in Hourly Earnings: An Analysis of Trends and Factors

Gender Pay Gap



In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px # is a high-level interface for creating various types of interactive
import plotly.graph_objects as go # is a lower-level interface that offers more control and customi
```

In [2]:

```
import warnings
warnings.filterwarnings('ignore')
```

In [3]:

```
df = pd.read_csv("female_hourly_earnings.csv")
```

In [4]:

```
df.shape
```

Out[4]:

```
(916, 7)
```

In [5]:

```
df.columns
```

Out[5]:

```
Index(['Unnamed: 0', 'country_id', 'country', 'gender_code', 'gender', 'year',  
      'amount_local_currency'],  
      dtype='object')
```

In [6]:

```
df = df.drop('Unnamed: 0', axis = 1)
```

In [7]:

```
df.duplicated().sum()
```

Out[7]:

```
0
```

In [8]:

```
df.isnull().sum()
```

Out[8]:

```
country_id      0  
country         0  
gender_code     0  
gender          0  
year           0  
amount_local_currency  0  
dtype: int64
```

In [9]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 916 entries, 0 to 915  
Data columns (total 6 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   country_id            916 non-null   int64  
1   country               916 non-null   object  
2   gender_code           916 non-null   object  
3   gender                916 non-null   object  
4   year                  916 non-null   int64  
5   amount_local_currency 916 non-null   float64  
dtypes: float64(1), int64(2), object(3)  
memory usage: 43.1+ KB
```

In [10]:

```
df.describe()
```

Out[10]:

	country_id	year	amount_local_currency
count	916.000000	916.000000	916.000000
mean	477.310044	2015.340611	350.170753
std	263.137563	3.331868	1502.006191
min	8.000000	2009.000000	-46.200000
25%	250.000000	2013.000000	9.227500
50%	498.000000	2015.000000	29.060000
75%	724.000000	2018.000000	57.627500
max	860.000000	2022.000000	20370.200000

In [11]:

```
object_columns = df.select_dtypes(include='object').columns.tolist()
numerical_columns = df.select_dtypes(include=['int', 'float']).columns.tolist()

print("Object columns:", object_columns)
print("Numerical columns:", numerical_columns)
```

```
Object columns: ['country', 'gender_code', 'gender']
Numerical columns: ['country_id', 'year', 'amount_local_currency']
```

In [12]:

```
df.nunique()
```

Out[12]:

```
country_id      45
country         45
gender_code      4
gender          4
year           14
amount_local_currency  829
dtype: int64
```

In [13]:

```
for i in object_columns:
    print(i)
    print(df[i].unique())
    print('\n')
```

country

```
['Armenia' 'Belgium' 'Germany' 'Bulgaria' 'Italy' 'Republic of Moldova'
'Denmark' 'United Kingdom' 'Slovakia' 'United States' 'Israel' 'Serbia'
'Iceland' 'Bosnia and Herzegovina' 'Croatia' 'Luxembourg' 'Estonia'
'Spain' 'Sweden' 'Latvia' 'Austria' 'Belarus' 'Albania' 'Türkiye' 'Malta'
'Russian Federation' 'Uzbekistan' 'Norway' 'Ukraine' 'North Macedonia'
'Poland' 'Hungary' 'Finland' 'Lithuania' 'Greece' 'Cyprus' 'Ireland'
'Montenegro' 'Slovenia' 'Czechia' 'Romania' 'Portugal' 'Switzerland'
'Netherlands' 'France']
```

gender_code

```
['F' 'M' 'T' 'G']
```

gender

```
['Female' 'Male' 'Total' 'Gender gap']
```

In [14]:

```
for i in object_columns:  
    print(i)  
    print(df[i].value_counts())  
    print('\n')
```

country	
United States	56
Switzerland	52
Republic of Moldova	52
France	40
Portugal	40
Türkiye	40
Spain	40
Czechia	36
Bosnia and Herzegovina	32
Greece	32
Sweden	32
Armenia	32
Serbia	32
Finland	28
Slovakia	24
Hungary	24
Belgium	24
United Kingdom	24
Iceland	20
Israel	16
Norway	12
Denmark	12
Lithuania	12
Slovenia	12
Italy	12
Malta	12
Bulgaria	12
Ireland	12
Germany	12
Austria	12
Latvia	12
Estonia	12
Luxembourg	12
Netherlands	12
Cyprus	12
Romania	8
Ukraine	8
Poland	8
Russian Federation	8
Albania	8
Montenegro	4
North Macedonia	4
Uzbekistan	4
Belarus	4
Croatia	4

Name: country, dtype: int64

gender_code	
F	229
M	229
T	229
G	229

Name: gender_code, dtype: int64


gender	
Female	229
Male	229
Total	229
Gender gap	229

Name: gender, dtype: int64


In [15]:

```
for i in object_columns:
    print('Countplot for:', i)
    plt.figure(figsize=(15,6))
    sns.countplot(df[i], data = df, palette = 'hls')
    plt.xticks(rotation = 90)
    plt.show()
    print('\n')
```


Countplot for: country



Countplot for: gender_code




Countplot for: gender



In [16]:

```
for i in object_columns:
    print('Pie plot for:', i)
    plt.figure(figsize=(20, 30))
    df[i].value_counts().plot(kind='pie', autopct='%1.1f%%')
    plt.title('Distribution of ' + i)
    plt.ylabel('')
    plt.show()
    print('\n')
```


Pie plot for: country



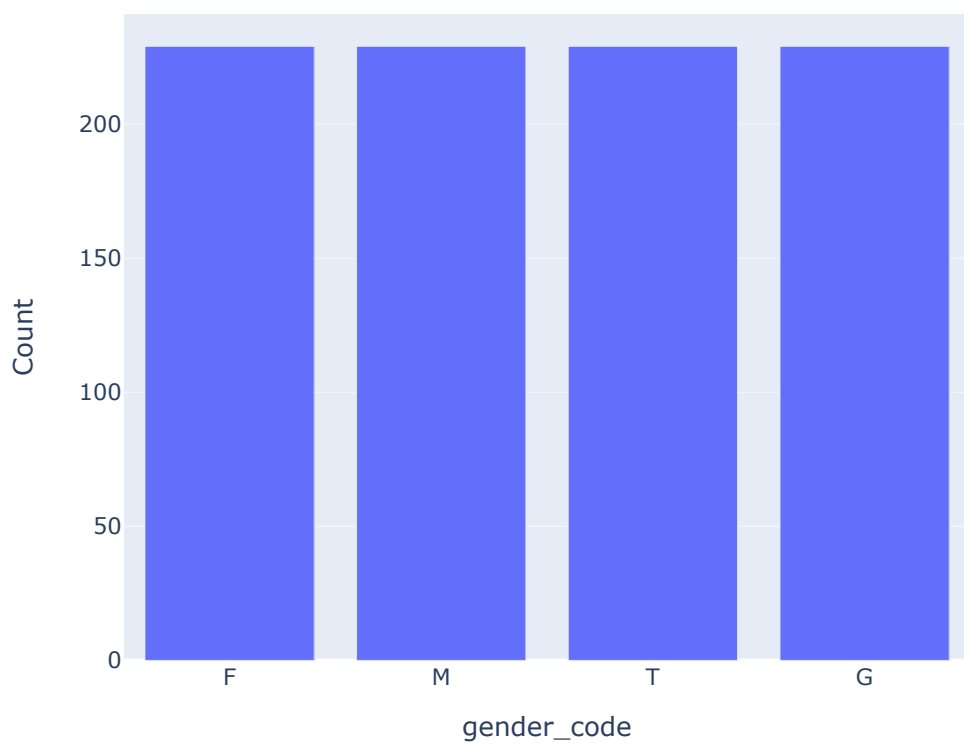
In [17]:

```
for i in object_columns:
    fig = go.Figure(data=[go.Bar(x=df[i].value_counts().index, y=df[i].value_counts()))
    fig.update_layout(
        title=i,
        xaxis_title=i,
        yaxis_title="Count")
    fig.show()
```


country



gender_code



gender




In [18]:

```
for i in object_columns:
    print('Pie plot for:', i)
    fig = px.pie(df, names=i, title='Distribution of ' + i)
    fig.show()
    print('\n')
```




Pie plot for: country

Distribution of country



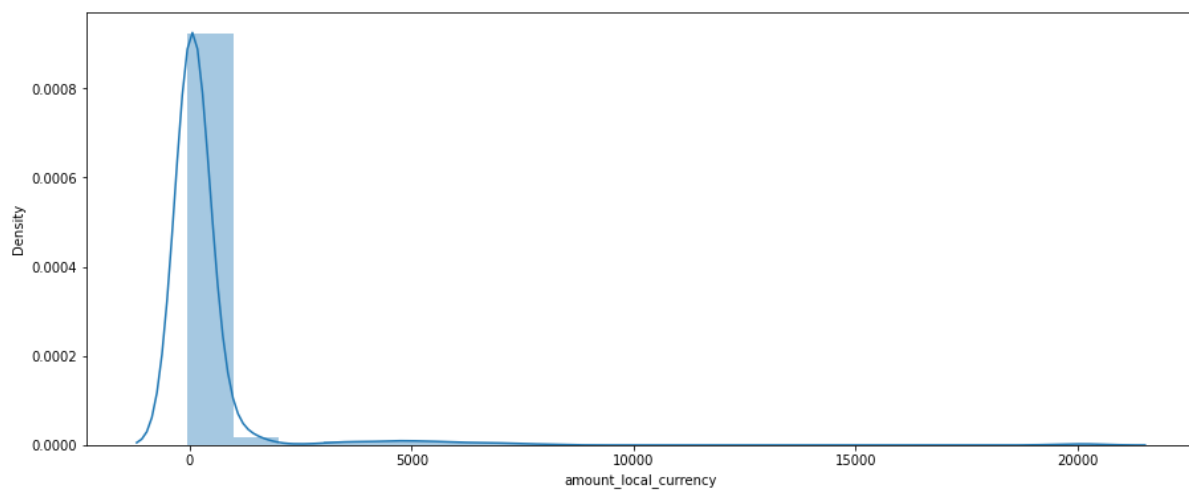


In [19]:

```
for i in numerical_columns:  
    plt.figure(figsize=(15,6))  
    sns.histplot(df[i], kde = True, bins = 20, palette = 'hls')  
    plt.xticks(rotation = 0)  
    plt.show()
```






In [20]:

```
for i in numerical_columns:  
    plt.figure(figsize=(15,6))  
    sns.distplot(df[i], kde = True, bins = 20)  
    plt.xticks(rotation = 0)  
    plt.show()
```






In [21]:

```
for i in numerical_columns:  
    plt.figure(figsize=(15,6))  
    sns.boxplot(df[i], data=df, palette='hls')  
    plt.xticks(rotation = 0)  
    plt.show()
```



In [22]:


```
for i in numerical_columns:  
    plt.figure(figsize=(15,6))  
    sns.violinplot(df[i], data=df, palette='hls')  
    plt.xticks(rotation = 0)  
    plt.show()
```




In [23]:

```
for i in numerical_columns:
    fig = go.Figure(data=[go.Histogram(x=df[i], nbinsx=20)])
    fig.update_layout(
        title=i,
        xaxis_title=i,
        yaxis_title="Count")
    fig.show()
```

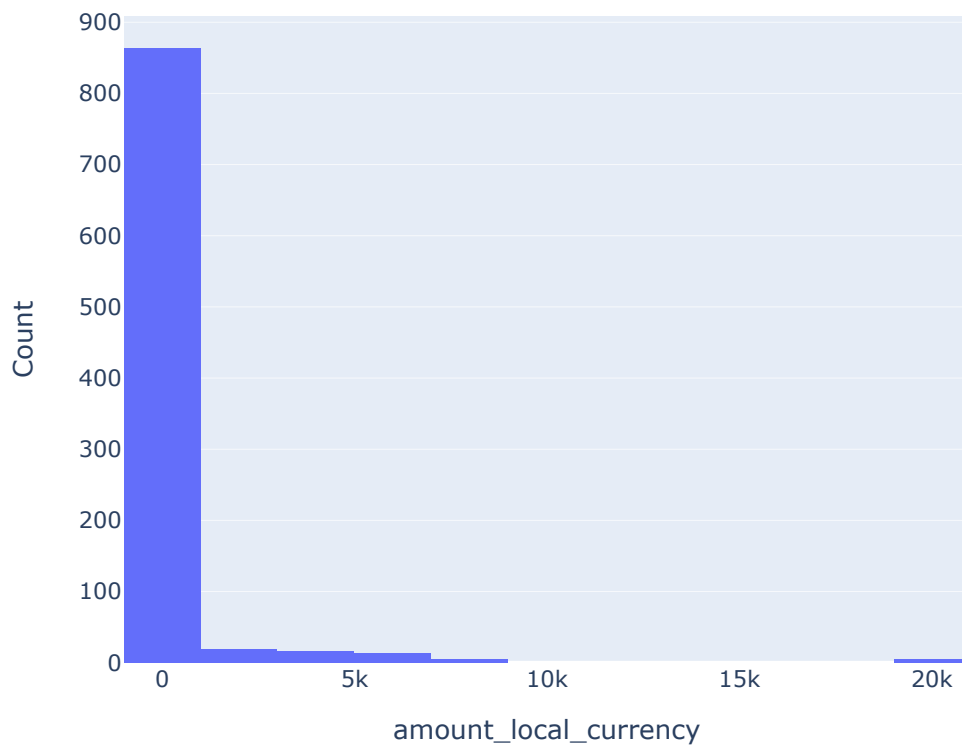
country_id



year



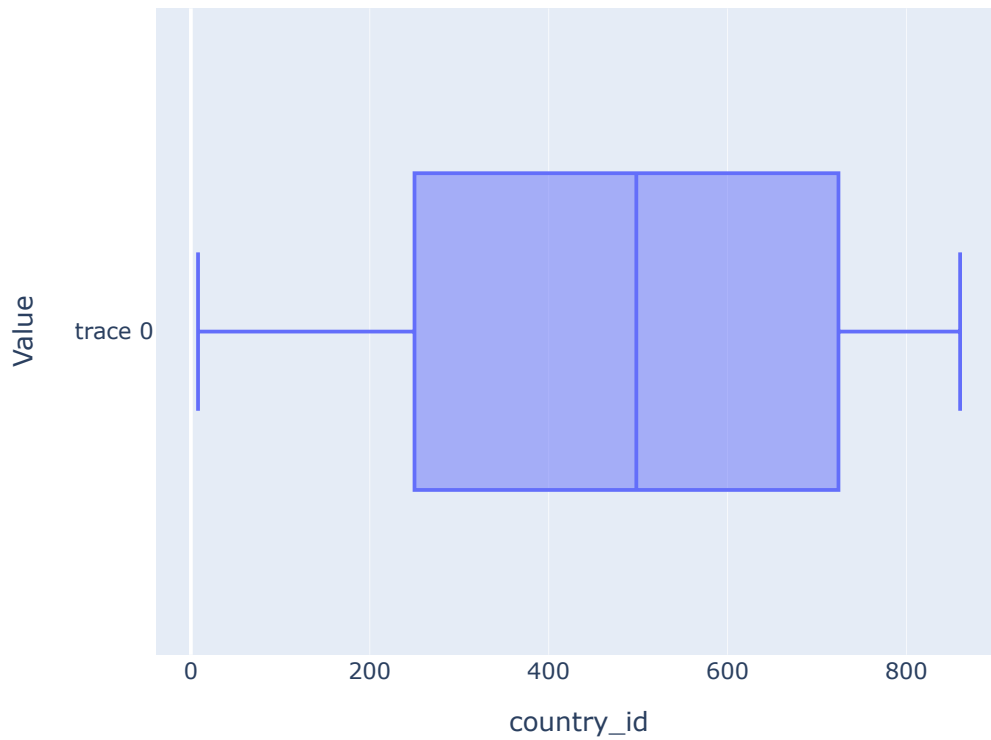
amount_local_currency




In [24]:

```
for i in numerical_columns:  
    fig = go.Figure(data=[go.Box(x=df[i])])  
    fig.update_layout(  
        title=i,  
        xaxis_title=i,  
        yaxis_title="Value")  
    fig.show()
```


country_id



year




amount_local_currency



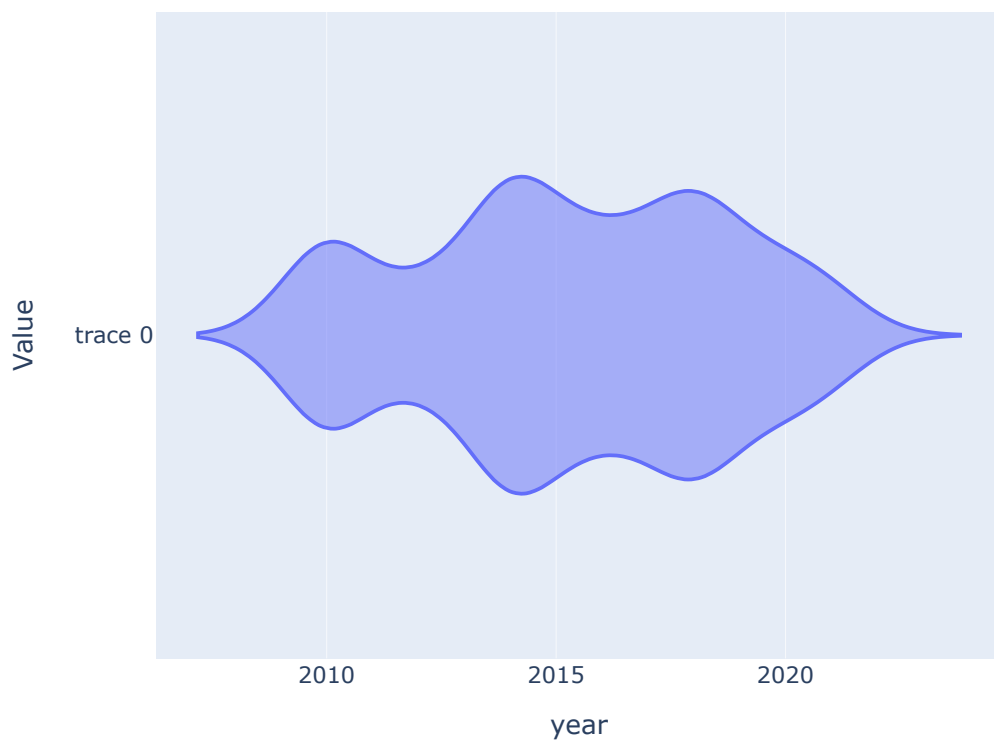
In [25]:

```
for i in numerical_columns:  
    fig = go.Figure(data=[go.Violin(x=df[i])])  
    fig.update_layout(  
        title=i,  
        xaxis_title=i,  
        yaxis_title="Value"  
    )  
    fig.show()
```

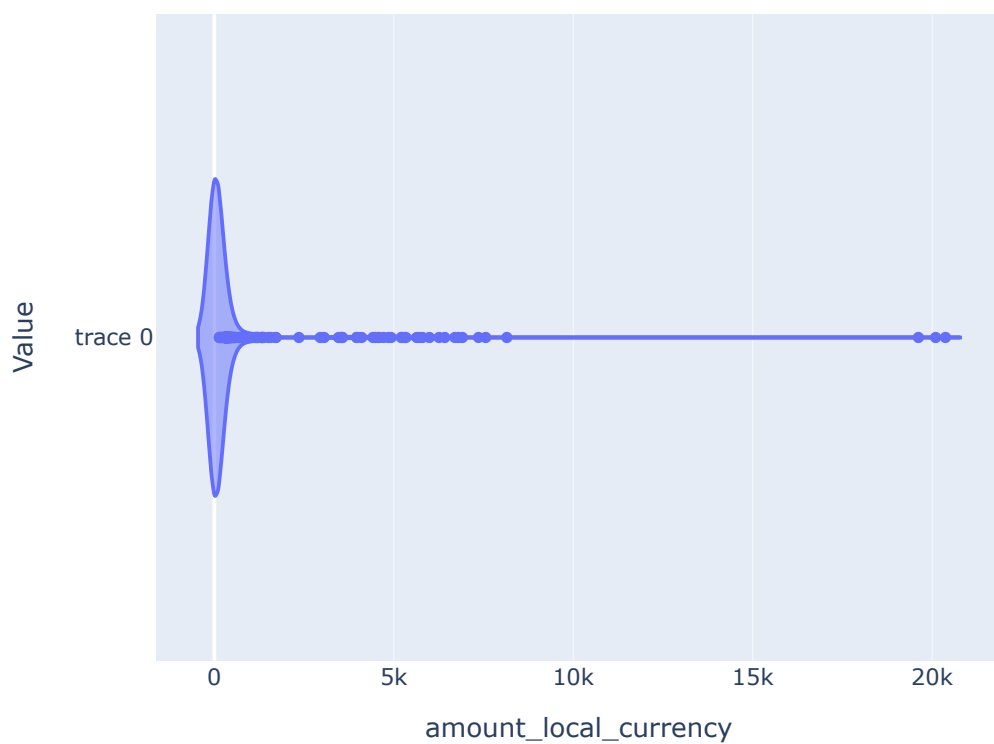
country_id



year




amount_local_currency






In [26]:

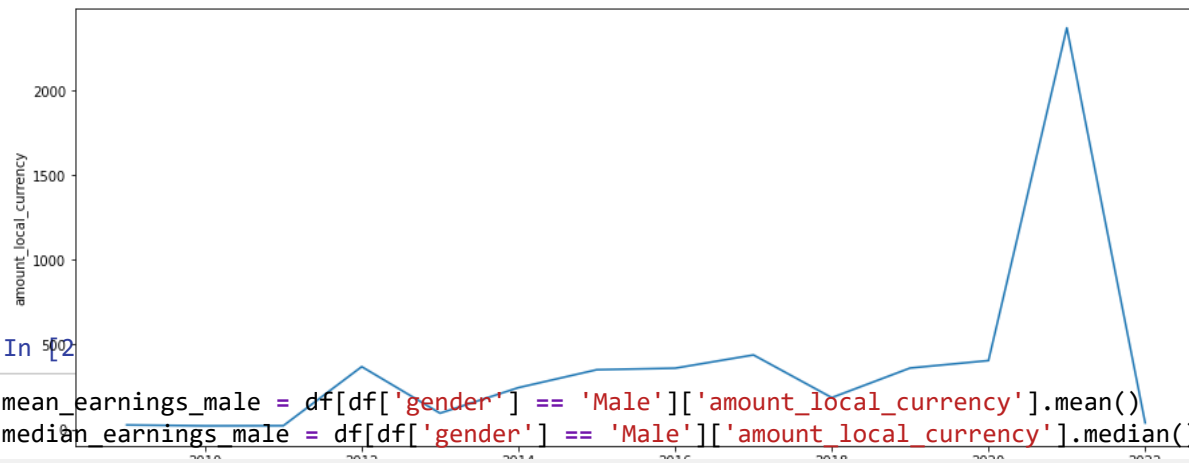


```
for i in numerical_columns:
    for j in object_columns:
        plt.figure(figsize=(15,6))
        sns.barplot(x = df[j], y = df[i], data = df, ci = None, palette = 'hls')
        plt.xticks(rotation = 90)
        plt.show()
```



In [27]:

```
for i in numerical_columns:
    for j in numerical_columns:
        if i != j:
            plt.figure(figsize=(15,6))
            sns.lineplot(x = df[j], y = df[i], data = df, ci = None, palette = 'hls')
            plt.show()
```





In [2]

```
mean_earnings_male = df[df['gender'] == 'Male']['amount_local_currency'].mean()
median_earnings_male = df[df['gender'] == 'Male']['amount_local_currency'].median()
```

In [29]:

```
mean_earnings_female = df[df['gender'] == 'Female']['amount_local_currency'].mean()
median_earnings_female = df[df['gender'] == 'Female']['amount_local_currency'].median()
```

In [30]:

```
print("Mean Earnings (Male):", mean_earnings_male)
print("Median Earnings (Male):", median_earnings_male)
print("Mean Earnings (Female):", mean_earnings_female)
print("Median Earnings (Female):", median_earnings_female)
```

Mean Earnings (Male): 478.3003930131005

Median Earnings (Male): 38.66

Mean Earnings (Female): 389.76903930131

Median Earnings (Female): 31.0

In [31]:

```
yearly_earnings = df.groupby(['year', 'gender'])['amount_local_currency'].mean().unstack()
```

In [32]:


```
yearly_earnings
```

Out[32]:

gender	Female	Gender gap	Male	Total
year				
2009	28.487500	5.712500	34.200000	32.125000
2010	21.190000	6.166429	27.356429	25.420714
2011	22.651667	5.926667	28.578333	26.475000
2012	367.716364	136.923636	504.640000	467.432727
2013	93.605455	36.262727	129.868182	118.183636
2014	253.120769	84.890256	338.011026	305.427949
2015	369.781053	113.021053	482.802105	439.235789
2016	386.647895	107.155263	493.803158	453.517895
2017	479.412353	120.770588	600.182941	553.461765
2018	203.024545	52.804242	255.828788	234.472424
2019	395.278667	100.317333	495.596000	454.092667
2020	441.477333	114.839333	556.316667	508.332000
2021	2877.045455	328.789091	3205.834545	3068.010000
2022	46.870000	3.820000	50.690000	48.990000

In [33]:


```
yearly_earnings.plot(kind='line')  
plt.xlabel('Year')  
plt.ylabel('Mean Earnings')  
plt.title('Mean Earnings Over Time by Gender')  
plt.show()
```



In [34]:

```
fig = px.line(yearly_earnings, x=yearly_earnings.index, y=['Male', 'Female'],
              title='Mean Earnings Over Time by Gender')
fig.update_xaxes(title='Year')
fig.update_yaxes(title='Mean Earnings')
fig.show()
```

Mean Earnings Over Time by Gender




In [35]:

```
earnings_by_year_gender = df.groupby(['year', 'gender'])['amount_local_currency'].mean().unstack()
```

In [36]:

```
plt.figure(figsize=(10, 6))
plt.plot(earnings_by_year_gender.index, earnings_by_year_gender['Male'], label='Male')
plt.plot(earnings_by_year_gender.index, earnings_by_year_gender['Female'], label='Female')
plt.xlabel('Year')
plt.ylabel('Mean Earnings')
plt.title('Mean Earnings Over Time by Gender')
plt.legend()
plt.grid(True)
plt.show()
```



In [37]:

```
gender_pay_gap = mean_earnings_male - mean_earnings_female
print("Gender Pay Gap:", gender_pay_gap)
```

Gender Pay Gap: 88.53135371179053

In [38]:

```
country_earnings = df.groupby(['country', 'gender'])['amount_local_currency'].mean().unstack()
```

In [39]:

```
country_earnings
```

Out[39]:

gender country	Female	Gender gap	Male	Total
Albania	2075.785000	667.200000	2742.985000	2566.040000
Armenia	831.622500	140.851250	972.473750	924.191250
Austria	26.943333	10.360000	37.303333	34.863333
Belarus	7.240000	2.740000	9.980000	8.460000
Belgium	33.866667	5.493333	39.360000	37.766667
Bosnia and Herzegovina	7.650000	0.692500	8.342500	8.178750
Bulgaria	9.153333	1.823333	10.976667	10.290000
Croatia	75.000000	31.000000	106.000000	92.000000
Cyprus	27.316667	3.463333	30.780000	29.993333
Czechia	303.611111	114.228889	417.840000	382.307778
Denmark	216.540000	58.783333	275.323333	257.520000
Estonia	9.270000	2.953333	12.223333	10.780000
Finland	33.975714	8.275714	42.251429	39.265714
France	19.231000	2.539000	21.770000	20.848000
Germany	32.916667	12.226667	45.143333	42.006667
Greece	10.491250	1.158750	11.650000	11.257500
Hungary	2870.078333	1032.506667	3902.585000	3456.608333
Iceland	5384.000000	1528.000000	6912.000000	6308.000000
Ireland	27.126667	8.676667	35.803333	32.006667
Israel	78.325000	15.900000	94.225000	89.350000
Italy	30.026667	13.093333	43.120000	40.200000
Latvia	7.680000	1.610000	9.290000	8.473333
Lithuania	6.420000	1.200000	7.620000	7.123333
Luxembourg	39.556667	11.636667	51.193333	48.890000
Malta	14.623333	4.160000	18.783333	17.396667
Montenegro	8.300000	0.530000	8.830000	8.680000
Netherlands	25.840000	7.653333	33.493333	31.293333
North Macedonia	302.120000	40.620000	342.740000	329.640000
Norway	247.400000	72.940000	320.340000	293.436667
Poland	23.220000	8.915000	32.135000	28.215000
Portugal	8.154000	1.954000	10.108000	9.394000
Republic of Moldova	29.875385	3.775385	33.650769	31.985385
Romania	17.785000	0.790000	18.575000	18.245000
Russian Federation	314.500000	135.000000	449.500000	385.000000
Serbia	371.841250	19.646250	391.487500	384.760000
Slovakia	11.220000	3.918333	15.138333	13.653333
Slovenia	18.140000	3.140000	21.280000	19.963333
Spain	26.991000	6.011000	33.002000	31.134000
Sweden	265.636250	42.762500	308.398750	291.610000

gender	Female	Gender gap	Male	Total
country				
Switzerland	51.126154	15.586154	66.712308	61.667692
Türkiye	26.958000	-1.256000	25.702000	25.958000
Ukraine	34.205000	8.980000	43.185000	38.840000
United Kingdom	21.915000	6.211667	28.126667	25.881667
United States	36.442857	7.015714	43.458571	40.382857
Uzbekistan	19616.000000	754.200000	20370.200000	20098.000000

In [40]:


```
country_earnings.plot(kind='bar', stacked=True)
plt.xlabel('Country')
plt.ylabel('Mean Earnings')
plt.title('Mean Earnings by Country and Gender')
plt.show()
```



In [41]:

```
fig = px.bar(country_earnings, x=country_earnings.index, y=['Male', 'Female'],
             title='Mean Earnings by Country and Gender')
fig.update_layout(width=800, height=500)
fig.update_xaxes(title_text='Country')
fig.update_yaxes(title_text='Mean Earnings')
fig.show()
```

Mean Earnings by Country and Gender



In [42]:

```
import statsmodels.api as sm
```

In [43]:

```
df_male = df[df['gender'] == 'Male']
df_female = df[df['gender'] == 'Female']
```

In [44]:

```
X_male = df_male[['year', 'country_id']]
X_male = sm.add_constant(X_male)
y_male = df_male['amount_local_currency']
```

In [45]:

```
model_male = sm.OLS(y_male, X_male).fit()
```

In [46]:

```
X_female = df_female[['year', 'country_id']]
X_female = sm.add_constant(X_female)
y_female = df_female['amount_local_currency']
```

In [47]:

```
model_female = sm.OLS(y_female, X_female).fit()
```

In [48]:

```
male_coefficients = model_male.params
female_coefficients = model_female.params
```

In [49]:

```
explained_gap = (male_coefficients - female_coefficients) @ X_female.mean()
unexplained_gap = (y_male.mean() - y_female.mean()) - explained_gap
```

In [50]:

```
print("Explained Gap:", explained_gap)
print("Unexplained Gap:", unexplained_gap)
```

```
Explained Gap: 88.53135371184908
Unexplained Gap: -5.854872142663226e-11
```

In [51]:

```
from scipy.stats import mannwhitneyu
```

In [52]:

```
gender_pay_gap_df = pd.DataFrame(columns=['Country', 'Year', 'p-value', 'Significant'])
```

In [53]:

```
for country in df['country'].unique():
    for year in df['year'].unique():
        data_country_year = df[(df['country'] == country) & (df['year'] == year)]
        male_earnings = data_country_year[data_country_year['gender'] == 'Male']['amount_local_currency']
        female_earnings = data_country_year[data_country_year['gender'] == 'Female']['amount_local_currency']
        if len(male_earnings) > 0 and len(female_earnings) > 0:
            stat, p_value = mannwhitneyu(male_earnings, female_earnings, alternative='two-sided')
            significance = 'Yes' if p_value < 0.05 else 'No'
            gender_pay_gap_df = gender_pay_gap_df.append({
                'Country': country,
                'Year': year,
                'p-value': p_value,
                'Significant': significance
            }, ignore_index=True)
```

In [54]:

```
print(gender_pay_gap_df)
```

	Country	Year	p-value	Significant
0	Armenia	2013	1.0	No
1	Armenia	2014	1.0	No
2	Armenia	2015	1.0	No
3	Armenia	2016	1.0	No
4	Armenia	2017	1.0	No
..
224	France	2018	1.0	No
225	France	2019	1.0	No
226	France	2020	1.0	No
227	France	2011	1.0	No
228	France	2012	1.0	No

[229 rows x 4 columns]